

# MLP-Mixer: An all-MLP Architecture for Vision

I. Tolstikhim, et al. Google Research, Brain Team  
Journal Club Discussion, REPU

January 31, 2022


# Index

- 1 Motivation
  - Key definitions
- 2 Principles
- 3 Approach
- 4 Results
- 5 Takeaways
  - Next Steps
- 6 References

# CNN and ViT are the go-to model for computer vision

- Convolution Neural Networks (CNN) are the current de-facto standard for computer vision applications (e.g. ResNet are the best solution for image classification).
- Vision Transformers (ViT) have reached a state-of-art performance and are based on self-attention layers (Google Research, Brain Team):
  - First success of Transformers in Natural Language Processing (NLP) [\[1\]](#).
  - Outperform by a short margin CNN's, with large enough dataset.

# SOTA applications in image recognition



**(a)** Mask R-CNN applications for object instance segmentation [2]




**(b)** Attention of learned positions, ViT [3]

# MLP-Mixer: a technically simple, but powerful solution

With modern hardware and larger available datasets, fully connected layer architectures such as MLP are working reasonably well [4]. This architecture is conceptually and technically simple alternative:

- Relies on simple matrix multiplication, reshapes, and transpositions operations.
- It does not follow a convolution or self-attention approach.



**Figure 2:** MLP basic concept that constitutes the MLP-Mixer [5]

# Key definitions

**Channel:** Number of colors that have the color model ( $RGB = 3$ ).

**Self-Attention:** Mechanism to weight neighboring instances to produce a contextual meaning.



**TPU-v3:** Google Cloud AI accelerator specifically for training NN.

**Optimization:** Seek to improve the accuracy of the model. Algorithms: SGD, Adam, Momentum, etc. Techniques: Batch-Normalization, Early Stopping, Gradient Clipping, etc.

# Each Mixer's layer use 2 MLP blocks connected

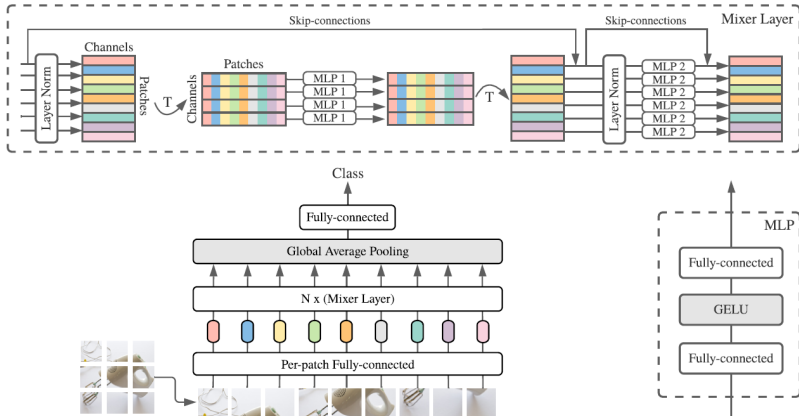
**Block 1:** Channel/Image-mixing MLP, and executes per-location operation. It shares parameters among layers.

**Block 2:** Token/Patches-mixing MLP, and executes a cross-location operation among channels.



**Figure 3:** (Left) Structure of each MLP Block. (Right) GELU benchmarking [6].

# MLP-Mixer Architecture is conceptually simple



**Figure 4:** Mixer Python implementation available at: <https://keras.io/examples/vision>



# Mixer is strongly influenced by CNNs and ViTs

**Table 1:** Similarities and differences between Mixer and modern vision architectures

	CNN	ViT
<i>Similarities</i>	<ul style="list-style-type: none"><li>- Parameter sharing</li><li>- Skip connections by addition (ResNets)</li><li>- Linear complexity according to image resolution</li></ul>	<ul style="list-style-type: none"><li>- Isotropic: same size input</li><li>- Convert images to tokens (patches)</li></ul>
<i>Differences</i>	<ul style="list-style-type: none"><li>- Pyramidal Structure with decreasing input size</li></ul>	<ul style="list-style-type: none"><li>- Positional Embeddings: Mixer is already sensible to input order.</li></ul>

# Mixer models were tested in various classification datasets

## Objective of the experiment

- ① Measure the **transfer accuracy** of performance for classification in popular datasets (ImageNet, CIFAR 100, etc.)
- ② Obtain test-time **throughput**: data units processed in a specific time.
- ③ Analyze the **computation cost** of pre-training the model and then fine tuning the model.

# Mixer's Experiments follow a transfer learning setup

## Pre-training

Resolution of 224 pixels using Adam, linear warmup of 10k steps, batch size 4096, weight decay, and gradient clipping.

## Fine-tuning

Use momentum, SGD, batch size 512, gradient clipping, cosine learning rate schedule, and increasing the input resolution.

## Metrics

Evaluate the trade-off between the model's computational cost and quality using: 1) Total pre-training time on TPUv3 accelerators for each training setup and 2) Throughput in images/sec/core.

## Models

Several Mixer models with different depths, patch resolutions, widths, input resolution, etc.

# Several model scales and combinations were defined


Specification	S/32	S/16	B/32	B/16	L/32	L/16	H/14
Number of layers	8	8	12	12	24	24	32
Patch resolution $P \times P$	$32 \times 32$	$16 \times 16$	$32 \times 32$	$16 \times 16$	$32 \times 32$	$16 \times 16$	$14 \times 14$
Hidden size $C$	512	512	768	768	1024	1024	1280
Sequence length $S$	49	196	49	196	49	196	256
MLP dimension $D_C$	2048	2048	3072	3072	4096	4096	5120
MLP dimension $D_S$	256	256	384	384	512	512	640
Parameters (M)	19	18	60	59	206	207	431

**Figure 5:** Specifications of the Mixer architectures, the "B", "L" and "H" mean base, large and huge model scale. The brief notation "B/16" means the model of the base scale with patches of resolution  $16 \times 16$  [5].

# Mixer has competitive accuracy and good trade-off


	ImNet top-1	Real top-1	Avg 5 top-1	VTAB-1k 19 tasks	Throughput img/sec/core	TPUv3 core-days
Pre-trained on ImageNet-21k (public)						
● HaloNet [51]	85.8	—	93.91	74.95	120	0.10k
● Mixer-L/16	84.15	87.86	94.39	72.72	105	0.41k
● ViT-L/16 [14]	85.30	88.62	94.04	70.64	32	0.18k
● BiT-R152x4 [22]	85.39	—	94.04	70.64	26	0.94k
Pre-trained on JFT-300M (proprietary)						
● NFNet-F4+ [7]	89.2	—	95.71	75.33	46	1.86k
● Mixer-H/14	87.94	90.18	95.33	76.29	40	1.01k
● BiT-R152x4 [22]	87.54	90.54	95.33	77.63	26	9.90k
● ViT-H/14 [14]	88.55	90.72	95.97	77.63	15	2.30k
Pre-trained on unlabelled or weakly labelled data (proprietary)						
● MPL [34]	90.0	91.12	—	—	—	20.48k
● ALIGN [21]	88.64	—	—	79.99	15	14.82k

**(a)** Big Mixer models has comparable metrics in most of the experiments



**(b)** Mixer performs in the Pareto frontier as well as extremely performant networks, in this case with ViT

# Model scale and training influence Mixer performance

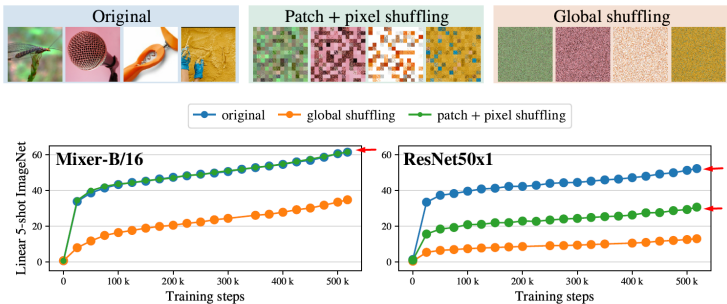


(a) Mixer is slightly below the frontier on the lower end of model scales, it sits confidently on the frontier at the high end

	Image size	Pre-Train Epochs	ImNet top-1	ResL top-1	Avg. S top-1 (img/sec/core)	Throughput top-1 (img/sec/core)	TPUv3 core-days
Pre-trained on ImageNet (with extra regularization)							
• Mixer-B/16	224	300	76.44	82.36	88.33	1384	0.01k <sup>(1)</sup>
• ViT-B/16 (■)	224	300	79.67	84.97	90.79	861	0.02k <sup>(1)</sup>
• Mixer-L/16	224	300	71.76	77.08	87.25	419	0.04k <sup>(1)</sup>
• ViT-L/16 (■)	224	300	76.11	80.93	89.66	280	0.05k <sup>(1)</sup>
Pre-trained on ImageNet-21k (with extra regularization)							
• Mixer-B/16	224	300	80.64	85.80	92.50	1384	0.15k <sup>(1)</sup>
• ViT-B/16 (■)	224	300	84.59	88.93	94.16	861	0.18k <sup>(1)</sup>
• Mixer-L/16	224	300	82.89	87.54	93.63	419	0.41k <sup>(1)</sup>
• ViT-L/16 (■)	224	300	84.46	88.35	94.49	280	0.55k <sup>(1)</sup>
• Mixer-L/16	448	300	83.91	87.75	93.86	105	0.41k <sup>(1)</sup>
Pre-trained on JFT-300M							
• Mixer-S/32	224	5	68.70	75.83	87.13	11489	0.01k
• Mixer-B/32	224	7	75.53	81.94	90.99	4208	0.05k
• Mixer-S/16	224	5	73.83	80.60	89.50	3994	0.03k
• BiT-R50x1	224	7	73.69	81.92	—	2159	0.08k
• Mixer-B/16	224	7	80.00	85.56	92.60	1384	0.08k
• Mixer-L/32	224	7	80.67	85.62	93.24	1314	0.12k
• BiT-R152x1	224	7	79.12	86.12	—	932	0.14k
• BiT-R50x2	224	7	78.92	86.06	—	890	0.14k
• BiT-R152x2	224	14	83.34	88.90	—	356	0.58k
• Mixer-L/16	224	7	84.05	88.14	94.51	419	0.23k
• Mixer-L/16	224	14	84.82	88.48	94.77	419	0.45k
• ViT-L/16	224	14	85.63	89.16	95.21	280	0.65k
• Mixer-H/14	224	14	86.32	89.14	95.49	194	1.01k
• BiT-R200x3	224	14	84.73	89.58	—	141	1.78k
• Mixer-L/16	448	14	86.78	89.72	95.13	105	0.45k
• ViT-H/14	224	14	86.65	89.56	95.57	87	2.35k
• ViT-L/16 [14]	512	14	87.76	90.54	95.63	32	0.65k

(b) As the pre-training dataset grows, Mixer's performance steadily improves.

# Model scale and training influence Mixer performance



(a) Mixer is invariant to the order of patches and pixels within the patches

# They describe a very simple architecture for vision

**Main Contribution:** In the era of powerful processors, Mixer demonstrates that it is as good as existing SOTA methods in terms of the *trade-off* between accuracy and computational resources required for training and inference.


- ✓ The next steps are to see whether such a design works in NLP or other domains.



# Mixer is the contextualization of my REPU project

Further analysis with a CKA kernel [7] for measure similarity among Mixer layer's activation might provide space for NN pruning techniques.

$$\text{CKA}(K, L) = \frac{\text{HSIC}(K, L)}{\sqrt{\text{HSIC}(K, K)\text{HSIC}(L, L)}}.$$



**Figure 9:** (Left) Definition of CKA kernel - Hilbert-Schmit Independence Criterion normalized. (Right) Some results running with a 8-layers Mixer. More updates in the Repo: <https://github.com/TheNewRobot/MixerCKA>.

# References I

- [1] A. Vaswani *et al.*, “Attention is all you need,” 2017.
- [2] K. He *et al.*, “Mask r-cnn,” 2018.
- [3] A. Dosovitskiy *et al.*, “An image is worth 16x16 words: Transformers for image recognition at scale,” 2021.
- [4] I. Goodfellow *et al.*, *Deep Learning*.  
MIT Press, 2016.  
<http://www.deeplearningbook.org>.
- [5] I. Tolstikhin *et al.*, “Mlp-mixer: An all-mlp architecture for vision,” 2021.

## References II

- [6] D. Hendrycks and K. Gimpel, “Gaussian error linear units (gelus),” 2020.
- [7] S. Kornblith, M. Norouzi, H. Lee, and G. Hinton, “Similarity of neural network representations revisited,” 2019.